


Support Vector Machine and Kernel Methods


Jiayu Zhou

¹Department of Computer Science and Engineering
Michigan State University
East Lansing, MI USA

Which Separator Do You Pick?




Robustness to Noisy Data



Being robust to noise (measurement error) is good (remember regularization).

Thicker Cushion Means More Robustness



We call such hyperplanes **fat**

Two Crucial Questions

- ❶ Can we efficiently find the fattest separating hyperplane?
- ❷ Is a fatter hyperplane better than a thin one?

Pulling Out the Bias

Before

$$\mathbf{x} \in \{1\} \times \mathbb{R}^d; \mathbf{w} \in \mathbb{R}^{d+1}$$

$$\mathbf{x} = \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_d \end{bmatrix}; \mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_d \end{bmatrix}$$

$$\text{signal} = \mathbf{w}^T \mathbf{x}$$

Pulling Out the Bias

Before

$$\mathbf{x} \in \{1\} \times \mathbb{R}^d; \mathbf{w} \in \mathbb{R}^{d+1}$$

$$\mathbf{x} = \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_d \end{bmatrix}; \mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_d \end{bmatrix}$$

$$\text{signal} = \mathbf{w}^T \mathbf{x}$$

After


$$\mathbf{x} \in \mathbb{R}^d; b \in \mathbb{R}, \mathbf{w} \in \mathbb{R}^d$$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_d \end{bmatrix}; \mathbf{w} = \begin{bmatrix} w_1 \\ \vdots \\ w_d \end{bmatrix}$$

bias b

$$\text{signal} = \mathbf{w}^T \mathbf{x} + b$$


Separating The Data



Hyperplane $h = (b, \mathbf{w})$
 h separates the data means:

$$y_n(\mathbf{w}^T \mathbf{x}_n + b) > 0$$

Separating The Data



Hyperplane $h = (b, \mathbf{w})$
 h separates the data means:


$$y_n(\mathbf{w}^T \mathbf{x}_n + b) > 0$$

By rescaling the weights and bias,

$$\min_{n=1, \dots, N} y_n(\mathbf{w}^T \mathbf{x}_n + b) = 1$$

Distance to the Hyperplane


- \mathbf{w} is normal to the hyperplane (why?)




Distance to the Hyperplane

- \mathbf{w} is normal to the hyperplane (why?)

$$\mathbf{w}^T(\mathbf{x}_2 - \mathbf{x}_1) = \mathbf{w}^T \mathbf{x}_2 - \mathbf{w}^T \mathbf{x}_1 = -b + b = 0$$



Distance to the Hyperplane



- \mathbf{w} is normal to the hyperplane (why?)


$$\mathbf{w}^T(\mathbf{x}_2 - \mathbf{x}_1) = \mathbf{w}^T \mathbf{x}_2 - \mathbf{w}^T \mathbf{x}_1 = -b + b = 0$$

- Scalar projection:

$$\mathbf{a}^T \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \cos(\mathbf{a}, \mathbf{b})$$

$$\Rightarrow \mathbf{a}^T \mathbf{b} / \|\mathbf{b}\| = \|\mathbf{a}\| \cos(\mathbf{a}, \mathbf{b})$$

Distance to the Hyperplane



- w is normal to the hyperplane (why?)

$$w^T(x_2 - x_1) = w^T x_2 - w^T x_1 = -b + b = 0$$


- Scalar projection:

$$a^T b = \|a\| \|b\| \cos(a, b)$$

$$\Rightarrow a^T b / \|b\| = \|a\| \cos(a, b)$$

- let x_\perp be the orthogonal projection of x to h , distance to hyperplane is given by projection of $x - x_\perp$ to w (why?)

Distance to the Hyperplane



- \mathbf{w} is normal to the hyperplane (why?)

$$\mathbf{w}^T(\mathbf{x}_2 - \mathbf{x}_1) = \mathbf{w}^T \mathbf{x}_2 - \mathbf{w}^T \mathbf{x}_1 = -b + b = 0$$

- Scalar projection:

$$\mathbf{a}^T \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \cos(\mathbf{a}, \mathbf{b})$$

$$\Rightarrow \mathbf{a}^T \mathbf{b} / \|\mathbf{b}\| = \|\mathbf{a}\| \cos(\mathbf{a}, \mathbf{b})$$

- let \mathbf{x}_\perp be the orthogonal projection of \mathbf{x} to h , distance to hyperplane is given by projection of $\mathbf{x} - \mathbf{x}_\perp$ to \mathbf{w} (why?)

$$\begin{aligned} \text{dist}(\mathbf{x}, h) &= \frac{1}{\|\mathbf{w}\|} \cdot |\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \mathbf{x}_\perp| \\ &= \frac{1}{\|\mathbf{w}\|} \cdot |\mathbf{w}^T \mathbf{x} + b| \end{aligned}$$

Fatness of a Separating Hyperplane


$$\text{dist}(\mathbf{x}, h) = \frac{1}{\|\mathbf{w}\|} \cdot |\mathbf{w}^T \mathbf{x} + b| = \frac{1}{\|\mathbf{w}\|} \cdot |y_n(\mathbf{w}^T \mathbf{x} + b)| = \frac{1}{\|\mathbf{w}\|} \cdot y_n(\mathbf{w}^T \mathbf{x} + b)$$

Fatness of a Separating Hyperplane

$$\text{dist}(\mathbf{x}, h) = \frac{1}{\|\mathbf{w}\|} \cdot |\mathbf{w}^T \mathbf{x} + b| = \frac{1}{\|\mathbf{w}\|} \cdot |y_n(\mathbf{w}^T \mathbf{x} + b)| = \frac{1}{\|\mathbf{w}\|} \cdot y_n(\mathbf{w}^T \mathbf{x} + b)$$

Fatness

= Distance to the closest point



$$\begin{aligned}\text{Fatness} &= \min_n \text{dist}(\mathbf{x}_n, h) \\ &= \frac{1}{\|\mathbf{w}\|} \min_n y_n(\mathbf{w}^T \mathbf{x} + b) \\ &= \frac{1}{\|\mathbf{w}\|}\end{aligned}$$

Maximizing the Margin

- Formal definition of margin:

$$\text{margin: } \gamma(h) = \frac{1}{\|\mathbf{w}\|}$$

Maximizing the Margin

- Formal definition of margin:

$$\text{margin: } \gamma(h) = \frac{1}{\|\mathbf{w}\|}$$

- NOTE: Bias b does not appear in the margin.

Maximizing the Margin

- Formal definition of margin:

$$\text{margin: } \gamma(h) = \frac{1}{\|\mathbf{w}\|}$$

- NOTE: Bias b does not appear in the margin.
- Objective maximizing margin:

$$\begin{aligned} \min_{b, \mathbf{w}} \quad & \frac{1}{2} \mathbf{w}^T \mathbf{w} \\ \text{subject to: } \quad & \min_{n=1, \dots, N} y_n (\mathbf{w}^T \mathbf{x}_n + b) = 1 \end{aligned}$$

Maximizing the Margin

- Formal definition of margin:

$$\text{margin: } \gamma(h) = \frac{1}{\|\mathbf{w}\|}$$

- NOTE: Bias b does not appear in the margin.
- Objective maximizing margin:

$$\begin{aligned} \min_{b, \mathbf{w}} \quad & \frac{1}{2} \mathbf{w}^T \mathbf{w} \\ \text{subject to:} \quad & \min_{n=1, \dots, N} y_n (\mathbf{w}^T \mathbf{x}_n + b) = 1 \end{aligned}$$

- An equivalent objective:

$$\begin{aligned} \min_{b, \mathbf{w}} \quad & \frac{1}{2} \mathbf{w}^T \mathbf{w} \\ \text{subject to:} \quad & y_n (\mathbf{w}^T \mathbf{x}_n + b) \geq 1 \text{ for } n = 1, \dots, N \end{aligned}$$

Example - Our Toy Data Set

$$\min_{b, \mathbf{w}} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

subject to: $y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1$ for $n = 1, \dots, N$

Training Data:

$$X = \begin{bmatrix} 0 & 0 \\ 2 & 2 \\ 2 & 0 \\ 3 & 0 \end{bmatrix}, \mathbf{y} = \begin{bmatrix} -1 \\ -1 \\ +1 \\ +1 \end{bmatrix}$$

What is the margin?

Example - Our Toy Data Set

$$\begin{aligned} \min_{b, \mathbf{w}} \quad & \frac{1}{2} \mathbf{w}^T \mathbf{w} \\ \text{subject to: } & y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 \text{ for } n = 1, \dots, N \end{aligned}$$

$$X = \begin{bmatrix} 0 & 0 \\ 2 & 2 \\ 2 & 0 \\ 3 & 0 \end{bmatrix}, \mathbf{y} = \begin{bmatrix} -1 \\ -1 \\ +1 \\ +1 \end{bmatrix} \Rightarrow \begin{cases} (1) : -b \geq 1 \\ (2) : -(2w_1 + 2w_2 + b) \geq 1 \\ (3) : 2w_1 + b \geq 1 \\ (4) : 3w_1 + b \geq 1 \end{cases}$$

$$\begin{cases} (1) + (3) & \rightarrow w_1 \geq 1 \\ (2) + (3) & \rightarrow w_2 \leq -1 \end{cases} \Rightarrow \frac{1}{2} \mathbf{w}^T \mathbf{w} = \frac{1}{2} (w_1^2 + w_2^2) \geq 1$$

Example - Our Toy Data Set

$$\begin{aligned} \min_{b, \mathbf{w}} \quad & \frac{1}{2} \mathbf{w}^T \mathbf{w} \\ \text{subject to: } & y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 \text{ for } n = 1, \dots, N \end{aligned}$$

$$X = \begin{bmatrix} 0 & 0 \\ 2 & 2 \\ 2 & 0 \\ 3 & 0 \end{bmatrix}, \mathbf{y} = \begin{bmatrix} -1 \\ -1 \\ +1 \\ +1 \end{bmatrix} \Rightarrow \begin{cases} (1) : -b \geq 1 \\ (2) : -(2w_1 + 2w_2 + b) \geq 1 \\ (3) : 2w_1 + b \geq 1 \\ (4) : 3w_1 + b \geq 1 \end{cases}$$

$$\begin{cases} (1) + (3) & \rightarrow w_1 \geq 1 \\ (2) + (3) & \rightarrow w_2 \leq -1 \end{cases} \Rightarrow \frac{1}{2} \mathbf{w}^T \mathbf{w} = \frac{1}{2} (w_1^2 + w_2^2) \geq 1$$

Thus: $w_1 = 1, w_2 = -1, b = -1$

Example - Our Toy Data Set

- Given data $X = \begin{bmatrix} 0 & 0 \\ 2 & 2 \\ 2 & 0 \\ 3 & 0 \end{bmatrix}$

- Optimal solution

$$\mathbf{w}^* = \begin{bmatrix} w_1 = 1 \\ w_2 = -1 \end{bmatrix}, b^* = -1$$

- Optimal hyperplane

$$g(\mathbf{x}) = \text{sign}(x_1 - x_2 - 1)$$


Example - Our Toy Data Set

- Given data $X = \begin{bmatrix} 0 & 0 \\ 2 & 2 \\ 2 & 0 \\ 3 & 0 \end{bmatrix}$

- Optimal solution

$$\mathbf{w}^* = \begin{bmatrix} w_1 = 1 \\ w_2 = -1 \end{bmatrix}, b^* = -1$$

- Optimal hyperplane
 $g(\mathbf{x}) = \text{sign}(x_1 - x_2 - 1)$
- margin:
 $\frac{1}{\|\mathbf{w}\|} = \frac{1}{\sqrt{2}} \approx 0.707$




Example - Our Toy Data Set

- Given data $X = \begin{bmatrix} 0 & 0 \\ 2 & 2 \\ 2 & 0 \\ 3 & 0 \end{bmatrix}$

- Optimal solution

$$\mathbf{w}^* = \begin{bmatrix} w_1 = 1 \\ w_2 = -1 \end{bmatrix}, b^* = -1$$

- Optimal hyperplane
 $g(\mathbf{x}) = \text{sign}(x_1 - x_2 - 1)$
- margin:
 $\frac{1}{\|\mathbf{w}\|} = \frac{1}{\sqrt{2}} \approx 0.707$




For data points (1), (2) and (3) $y_n(\mathbf{x}_n^T \mathbf{w}^* + b^*) = 1$

Example - Our Toy Data Set

- Given data $X = \begin{bmatrix} 0 & 0 \\ 2 & 2 \\ 2 & 0 \\ 3 & 0 \end{bmatrix}$
- Optimal solution

$$\mathbf{w}^* = \begin{bmatrix} w_1 = 1 \\ w_2 = -1 \end{bmatrix}, b^* = -1$$

- Optimal hyperplane
 $g(\mathbf{x}) = \text{sign}(x_1 - x_2 - 1)$
- margin:
 $\frac{1}{\|\mathbf{w}\|} = \frac{1}{\sqrt{2}} \approx 0.707$



For data points (1), (2) and
(3) $y_n(\mathbf{x}_n^T \mathbf{w}^* + b^*) = 1$
Support Vectors

Solver: Quadratic Programming

$$\min_{\mathbf{u} \in \mathbb{R}^q} \quad \frac{1}{2} \mathbf{u}^T Q \mathbf{u} + \mathbf{p}^T \mathbf{u}$$

subject to: $A\mathbf{u} \geq \mathbf{c}$

$$\mathbf{u}^* \leftarrow QP(Q, \mathbf{p}, A, \mathbf{c})$$

($Q = 0$ is linear programming.)

<http://cvxopt.org/examples/tutorial/qp.html>

Maximum Margin Hyperplane is QP

$$\min_{b, \mathbf{w}} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

subject to: $y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1, \forall n$

$$\mathbf{u} = \begin{bmatrix} b \\ \mathbf{w} \end{bmatrix} \in \mathbb{R}^{d+1} \Rightarrow \frac{1}{2} \mathbf{w}^T \mathbf{w} = [b, \mathbf{w}^T] \begin{bmatrix} 0 & \mathbf{0}_d^T \\ \mathbf{0}_d & I_d \end{bmatrix} \begin{bmatrix} b \\ \mathbf{w}^T \end{bmatrix} = \mathbf{u}^T \begin{bmatrix} 0 & \mathbf{0}_d^T \\ \mathbf{0}_d & I_d \end{bmatrix} \mathbf{u}$$

$$\min_{\mathbf{u} \in \mathbb{R}^q} \quad \frac{1}{2} \mathbf{u}^T Q \mathbf{u} + \mathbf{p}^T \mathbf{u}$$

subject to: $A \mathbf{u} \geq \mathbf{c}$

Maximum Margin Hyperplane is QP

$$\min_{b, \mathbf{w}} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

subject to: $y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1, \forall n$

$$\min_{\mathbf{u} \in \mathbb{R}^q} \quad \frac{1}{2} \mathbf{u}^T Q \mathbf{u} + \mathbf{p}^T \mathbf{u}$$

subject to: $A \mathbf{u} \geq \mathbf{c}$

$$\mathbf{u} = \begin{bmatrix} b \\ \mathbf{w} \end{bmatrix} \in \mathbb{R}^{d+1} \Rightarrow \frac{1}{2} \mathbf{w}^T \mathbf{w} = [b, \mathbf{w}^T] \begin{bmatrix} 0 & \mathbf{0}_d^T \\ \mathbf{0}_d & I_d \end{bmatrix} \begin{bmatrix} b \\ \mathbf{w}^T \end{bmatrix} = \mathbf{u}^T \begin{bmatrix} 0 & \mathbf{0}_d^T \\ \mathbf{0}_d & I_d \end{bmatrix} \mathbf{u}$$

$$Q = \begin{bmatrix} 0 & \mathbf{0}_d^T \\ \mathbf{0}_d & I_d \end{bmatrix}, \mathbf{p} = \mathbf{0}_{d+1}$$

Maximum Margin Hyperplane is QP

$$\min_{b, \mathbf{w}} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

$$\text{subject to: } y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1, \forall n$$

$$\min_{\mathbf{u} \in \mathbb{R}^q} \quad \frac{1}{2} \mathbf{u}^T Q \mathbf{u} + \mathbf{p}^T \mathbf{u}$$

$$\text{subject to: } A \mathbf{u} \geq \mathbf{c}$$

$$\mathbf{u} = \begin{bmatrix} b \\ \mathbf{w} \end{bmatrix} \in \mathbb{R}^{d+1} \Rightarrow \frac{1}{2} \mathbf{w}^T \mathbf{w} = [b, \mathbf{w}^T] \begin{bmatrix} 0 & \mathbf{0}_d^T \\ \mathbf{0}_d & I_d \end{bmatrix} \begin{bmatrix} b \\ \mathbf{w}^T \end{bmatrix} = \mathbf{u}^T \begin{bmatrix} 0 & \mathbf{0}_d^T \\ \mathbf{0}_d & I_d \end{bmatrix} \mathbf{u}$$

$$Q = \begin{bmatrix} 0 & \mathbf{0}_d^T \\ \mathbf{0}_d & I_d \end{bmatrix}, \mathbf{p} = \mathbf{0}_{d+1}$$

$$y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 = [y_n, y_n \mathbf{x}_n^T] \mathbf{u} \geq 1 \Rightarrow \begin{bmatrix} y_1 & y_1 \mathbf{x}_1^T \\ \vdots & \vdots \\ y_N & y_N \mathbf{x}_N^T \end{bmatrix} \mathbf{u} \geq \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix}$$

Maximum Margin Hyperplane is QP

$$\min_{b, \mathbf{w}} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

$$\text{subject to: } y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1, \forall n$$

$$\min_{\mathbf{u} \in \mathbb{R}^q} \quad \frac{1}{2} \mathbf{u}^T Q \mathbf{u} + \mathbf{p}^T \mathbf{u}$$

$$\text{subject to: } A \mathbf{u} \geq \mathbf{c}$$

$$\mathbf{u} = \begin{bmatrix} b \\ \mathbf{w} \end{bmatrix} \in \mathbb{R}^{d+1} \Rightarrow \frac{1}{2} \mathbf{w}^T \mathbf{w} = [b, \mathbf{w}^T] \begin{bmatrix} 0 & \mathbf{0}_d^T \\ \mathbf{0}_d & I_d \end{bmatrix} \begin{bmatrix} b \\ \mathbf{w}^T \end{bmatrix} = \mathbf{u}^T \begin{bmatrix} 0 & \mathbf{0}_d^T \\ \mathbf{0}_d & I_d \end{bmatrix} \mathbf{u}$$

$$Q = \begin{bmatrix} 0 & \mathbf{0}_d^T \\ \mathbf{0}_d & I_d \end{bmatrix}, \mathbf{p} = \mathbf{0}_{d+1}$$

$$y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 = [y_n, y_n \mathbf{x}_n^T] \mathbf{u} \geq 1 \Rightarrow \begin{bmatrix} y_1 & y_1 \mathbf{x}_1^T \\ \vdots & \vdots \\ y_N & y_N \mathbf{x}_N^T \end{bmatrix} \mathbf{u} \geq \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix}$$

$$A = \begin{bmatrix} y_1 & y_1 \mathbf{x}_1^T \\ \vdots & \vdots \\ y_N & y_N \mathbf{x}_N^T \end{bmatrix}, \mathbf{c} = \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix}$$

Back To Our Example

Exercise:

$$X = \begin{bmatrix} 0 & 0 \\ 2 & 2 \\ 2 & 0 \\ 3 & 0 \end{bmatrix}, \mathbf{y} = \begin{bmatrix} -1 \\ -1 \\ +1 \\ +1 \end{bmatrix} \quad \left\{ \begin{array}{l} (1) : -b \geq 1 \\ (2) : -(2w_1 + 2w_2 + b) \geq 1 \\ (3) : 2w_1 + b \geq 1 \\ (4) : 3w_1 + b \geq 1 \end{array} \right.$$

Show the corresponding $Q, \mathbf{p}, A, \mathbf{c}$.

Back To Our Example

Exercise:

$$X = \begin{bmatrix} 0 & 0 \\ 2 & 2 \\ 2 & 0 \\ 3 & 0 \end{bmatrix}, \mathbf{y} = \begin{bmatrix} -1 \\ -1 \\ +1 \\ +1 \end{bmatrix} \quad \left\{ \begin{array}{l} (1) : -b \geq 1 \\ (2) : -(2w_1 + 2w_2 + b) \geq 1 \\ (3) : 2w_1 + b \geq 1 \\ (4) : 3w_1 + b \geq 1 \end{array} \right.$$

Show the corresponding $Q, \mathbf{p}, A, \mathbf{c}$.

$$Q = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}, \mathbf{p} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}, A = \begin{bmatrix} -1 & 0 & 0 \\ -1 & -2 & -2 \\ 1 & 2 & 0 \\ 1 & 3 & 0 \end{bmatrix}, \mathbf{c} = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \end{bmatrix}$$

Back To Our Example

Exercise:

$$X = \begin{bmatrix} 0 & 0 \\ 2 & 2 \\ 2 & 0 \\ 3 & 0 \end{bmatrix}, \mathbf{y} = \begin{bmatrix} -1 \\ -1 \\ +1 \\ +1 \end{bmatrix} \quad \begin{cases} (1) : -b \geq 1 \\ (2) : -(2w_1 + 2w_2 + b) \geq 1 \\ (3) : 2w_1 + b \geq 1 \\ (4) : 3w_1 + b \geq 1 \end{cases}$$

Show the corresponding $Q, \mathbf{p}, A, \mathbf{c}$.

$$Q = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}, \mathbf{p} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}, A = \begin{bmatrix} -1 & 0 & 0 \\ -1 & -2 & -2 \\ 1 & 2 & 0 \\ 1 & 3 & 0 \end{bmatrix}, \mathbf{c} = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \end{bmatrix}$$

Use your QP-solver to give

$$\mathbf{u}^* = [b^*, w_1^*, w_2^*]^T = [-1, 1, -1]$$

Primal QP algorithm for linear-SVM

- ❶ Let $\mathbf{p} = \mathbf{0}_{d+1}$ be the $(d+1)$ -vector of zeros and $\mathbf{c} = \mathbf{1}_N$ the N -vector of ones. Construct matrices Q and A , where

$$A = \begin{bmatrix} y_1 & -y_1 \mathbf{x}_1^T \\ \vdots & \vdots \\ y_N & -y_N \mathbf{x}_N^T \end{bmatrix}, Q = \begin{bmatrix} 0 & \mathbf{0}_d^T \\ \mathbf{0}_d & I_d \end{bmatrix}$$


- ❷ Return $\begin{bmatrix} b^* \\ \mathbf{w}^* \end{bmatrix} = \mathbf{u}^* \leftarrow QP(Q, \mathbf{p}, A, \mathbf{c})$.
- ❸ The final hypothesis is $g(\mathbf{x}) = \text{sign}(\mathbf{x}^T \mathbf{w}^* + b^*)$.

Link to Regularization


$$\begin{aligned} \min_{\mathbf{w}} \quad & E_{in}(\mathbf{w}) \\ \text{subject to: } & \mathbf{w}^T \mathbf{w} \leq C \end{aligned}$$

	optimal hyperplane	regularization
minimize	$\mathbf{w}^T \mathbf{w}$	E_{in}
subject to	$E_{in} = 0$	$\mathbf{w}^T \mathbf{w} \leq C$

How to Handle Non-Separable Data?




(a) Few noisy data.




(b) Nonlinearly separable.

How to Handle Non-Separable Data?



(a) Few noisy data.



(b) Nonlinearly separable.

- (a) Tolerate noisy data points: soft-margin SVM.
- (b) Inherent nonlinear boundary: non-linear transformation.

Non-Linear Transformation

$$\Phi_1(\mathbf{x}) = (x_1, x_2)$$

$$\Phi_2(\mathbf{x}) = (x_1, x_2, x_1^2, x_1x_2, x_2^2)$$

$$\Phi_3(\mathbf{x}) = (x_1, x_2, x_1^2, x_1x_2, x_2^2, x_1^3, x_1^2x_2, x_1x_2^2, x_2^3)$$

Non-Linear Transformation

- Using the nonlinear transform with the optimal hyperplane using a transform $\Phi: \mathbb{R}^d \rightarrow \mathbb{R}^{\tilde{d}}$:

$$\mathbf{z}_n = \Phi(\mathbf{x}_n)$$

Non-Linear Transformation

- Using the nonlinear transform with the optimal hyperplane using a transform $\Phi: \mathbb{R}^d \rightarrow \mathbb{R}^{\tilde{d}}$:

$$\mathbf{z}_n = \Phi(\mathbf{x}_n)$$

- Solve the hard-margin SVM in the \mathcal{Z} -space $(\tilde{\mathbf{w}}^*, \tilde{b}^*)$:

$$\min_{\tilde{b}, \tilde{\mathbf{w}}} \frac{1}{2} \tilde{\mathbf{w}}^T \tilde{\mathbf{w}}$$

$$\text{subject to: } y_n(\tilde{\mathbf{w}}^T \mathbf{z}_n + \tilde{b}) \geq 1, \forall n$$

Non-Linear Transformation

- Using the nonlinear transform with the optimal hyperplane using a transform $\Phi: \mathbb{R}^d \rightarrow \mathbb{R}^{\tilde{d}}$:

$$\mathbf{z}_n = \Phi(\mathbf{x}_n)$$

- Solve the hard-margin SVM in the \mathcal{Z} -space $(\tilde{\mathbf{w}}^*, \tilde{b}^*)$:


$$\min_{\tilde{b}, \tilde{\mathbf{w}}} \frac{1}{2} \tilde{\mathbf{w}}^T \tilde{\mathbf{w}}$$

$$\text{subject to: } y_n(\tilde{\mathbf{w}}^T \mathbf{z}_n + \tilde{b}) \geq 1, \forall n$$


- Final hypothesis:

$$g(\mathbf{x}) = \text{sign}(\tilde{\mathbf{w}}^{*T} \Phi(\mathbf{x}) + \tilde{b}^*)$$

SVM and non-linear transformation




(a) SVM with Φ_2




(b) SVM with Φ_3

The margin is shaded in yellow, and the support vectors are boxed.

SVM and non-linear transformation



(a) SVM with Φ_2




(b) SVM with Φ_3


The margin is shaded in yellow, and the support vectors are boxed.

- For Φ_2 , $\tilde{d}_2 = 5$ and for Φ_3 , $\tilde{d}_3 = 9$

SVM and non-linear transformation



(a) SVM with Φ_2



(b) SVM with Φ_3

The margin is shaded in yellow, and the support vectors are boxed.

- For Φ_2 , $\tilde{d}_2 = 5$ and for Φ_3 , $\tilde{d}_3 = 9$
- \tilde{d}_3 is nearly double \tilde{d}_2 , yet the resulting SVM separator is not severely overfitting with Φ_3 (regularization?).

Support Vector Machine Summary

- A very powerful, easy to use linear model which comes with **automatic** regularization.

Support Vector Machine Summary

- A very powerful, easy to use linear model which comes with **automatic** regularization.
- Fully exploit SVM: Kernel
 - potential robustness to overfitting even after transforming to a much higher dimension
 - How about infinite dimensional transforms?
 - **Kernel Trick**

SVM Dual: Formulation

- Primal and dual in optimization.

SVM Dual: Formulation

- Primal and dual in optimization.
- The dual view of SVM enables us to exploit the kernel trick.

SVM Dual: Formulation

- Primal and dual in optimization.
- The dual view of SVM enables us to exploit the kernel trick.
- In the primal SVM problem we solve $\mathbf{w} \in \mathbb{R}^d, b$, while in the dual problem we solve $\alpha \in \mathbb{R}^N$

$$\begin{aligned} \max_{\alpha \in \mathbb{R}^N} \quad & \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{m=1}^N \sum_{n=1}^N \alpha_n \alpha_m y_n y_m \mathbf{x}_n^T \mathbf{x}_m \\ \text{subject to} \quad & \sum_{n=1}^N y_n \alpha_n = 0, \alpha_n \geq 0, \forall n \end{aligned}$$

which is also a QP problem.

SVM Dual: Prediction

- We can obtain the primal solution:

$$\mathbf{w}^* = \sum_{n=1}^N y_n \alpha_n^* \mathbf{x}_n$$

where for support vectors $\alpha_n > 0$

SVM Dual: Prediction

- We can obtain the primal solution:

$$\mathbf{w}^* = \sum_{n=1}^N y_n \alpha_n^* \mathbf{x}_n$$

where for support vectors $\alpha_n > 0$

- The optimal hypothesis:

$$\begin{aligned} g(\mathbf{x}) &= \text{sign}(\mathbf{w}^{*T} \mathbf{x} + b^*) \\ &= \text{sign} \left(\sum_{n=1}^N y_n \alpha_n^* \mathbf{x}_n^T \mathbf{x} + b^* \right) \\ &= \text{sign} \left(\sum_{\alpha_n^* > 0} y_n \alpha_n^* \mathbf{x}_n^T \mathbf{x} + b^* \right) \end{aligned}$$

Dual SVM: Summary

$$\begin{aligned} & \max_{\boldsymbol{\alpha} \in \mathbb{R}^N} \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{m=1}^N \sum_{n=1}^N \alpha_n \alpha_m y_n y_m \mathbf{x}_n^T \mathbf{x}_m \\ & \text{subject to } \sum_{n=1}^N y_n \alpha_n = 0, \alpha_n \geq 0, \forall n \end{aligned}$$

$$\mathbf{w}^* = \sum_{n=1}^N y_n \alpha_n^* \mathbf{x}_n$$

Common SVM Basis Functions

- \mathbf{z}_k = polynomial terms of \mathbf{x}_k of degree 1 to q
- \mathbf{z}_k = radial basis function of \mathbf{x}_k

$$\mathbf{z}_k(j) = \phi_j(\mathbf{x}_k) = \exp(-|\mathbf{x}_k - \mathbf{c}_j|^2/\sigma^2)$$

- \mathbf{z}_k = sigmoid functions of \mathbf{x}_k

Quadratic Basis Functions

$$\Phi(\mathbf{x}) = \begin{bmatrix} 1 \\ \sqrt{2}x_1 \\ \vdots \\ \sqrt{2}x_d \\ x_1^2 \\ \vdots \\ x_d^2 \\ \sqrt{2}x_1x_2 \\ \vdots \\ \sqrt{2}x_1x_d \\ \sqrt{2}x_2x_3 \\ \vdots \\ \sqrt{2}x_{d-1}x_d \end{bmatrix}$$

- Including **Constant Term**, **Linear Terms**, **Pure Quadratic Terms**, **Quadratic Cross-Terms**
- The number of terms is approximately $d^2/2$.
- You may be wondering what those $\sqrt{2}$ s are doing. You'll find out why they're there soon.

Dual SVM: Non-linear Transformation

$$\begin{aligned} \max_{\alpha \in \mathbb{R}^N} \quad & \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{m=1}^N \sum_{n=1}^N \alpha_n \alpha_m y_n y_m \Phi(\mathbf{x}_n)^T \Phi(\mathbf{x}_m) \\ \text{subject to} \quad & \sum_{n=1}^N y_n \alpha_n = 0, \alpha_n \geq 0, \forall n \end{aligned}$$

$$\mathbf{w}^* = \sum_{n=1}^N y_n \alpha_n^* \Phi(\mathbf{x}_n)$$

- Need to prepare a matrix Q , $Q_{nm} = y_n y_m \Phi(\mathbf{x}_n)^T \Phi(\mathbf{x}_m)$
- Cost?

Dual SVM: Non-linear Transformation

$$\begin{aligned} \max_{\alpha \in \mathbb{R}^N} \quad & \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{m=1}^N \sum_{n=1}^N \alpha_n \alpha_m y_n y_m \Phi(\mathbf{x}_n)^T \Phi(\mathbf{x}_m) \\ \text{subject to} \quad & \sum_{n=1}^N y_n \alpha_n = 0, \alpha_n \geq 0, \forall n \end{aligned}$$

$$\mathbf{w}^* = \sum_{n=1}^N y_n \alpha_n^* \Phi(\mathbf{x}_n)$$

- Need to prepare a matrix Q , $Q_{nm} = y_n y_m \Phi(\mathbf{x}_n)^T \Phi(\mathbf{x}_m)$
- Cost?
 - We must do $N^2/2$ dot products to get this matrix ready.
 - Each dot product requires $d^2/2$ additions and multiplications, The whole thing costs $N^2 d^2/4$.

Quadratic Dot Products

$$\Phi(\mathbf{a})^T \Phi(\mathbf{b}) = \begin{bmatrix} 1 \\ \sqrt{2}a_1 \\ \vdots \\ \sqrt{2}a_m \\ a_1^2 \\ \vdots \\ a_m^2 \\ \sqrt{2}a_1a_2 \\ \vdots \\ \sqrt{2}a_1a_d \\ \sqrt{2}a_2a_3 \\ \vdots \\ \sqrt{2}a_{d-1}a_d \end{bmatrix}^T \begin{bmatrix} 1 \\ \sqrt{2}b_1 \\ \vdots \\ \sqrt{2}b_d \\ b_1^2 \\ \vdots \\ b_d^2 \\ \sqrt{2}b_1b_2 \\ \vdots \\ \sqrt{2}b_1b_d \\ \sqrt{2}b_2b_3 \\ \vdots \\ \sqrt{2}b_{d-1}b_d \end{bmatrix}$$

- Constant Term 1
- Linear Terms

$$\sum_{i=1}^d 2a_i b_i$$

- Pure Quadratic Terms

$$\sum_{i=1}^d a_i^2 b_i^2$$

- Quadratic Cross-Terms

$$\sum_{i=1}^d \sum_{j=i+1}^d 2a_i a_j b_i b_j$$

Quadratic Dot Product

- Does $\Phi(\mathbf{a})^T \Phi(\mathbf{b})$ look familiar?

$$\Phi(\mathbf{a})^T \Phi(\mathbf{b}) = 1 + 2 \sum_{i=1}^d a_i b_i + \sum_{i=1}^d a_i^2 b_i^2 + \sum_{i=1}^d \sum_{j=i+1}^d 2a_i a_j b_i b_j$$

Quadratic Dot Product

- Does $\Phi(\mathbf{a})^T \Phi(\mathbf{b})$ look familiar?

$$\Phi(\mathbf{a})^T \Phi(\mathbf{b}) = 1 + 2 \sum_{i=1}^d a_i b_i + \sum_{i=1}^d a_i^2 b_i^2 + \sum_{i=1}^d \sum_{j=i+1}^d 2a_i a_j b_i b_j$$

- Try this: $(\mathbf{a}^T \mathbf{b} + 1)^2$

Quadratic Dot Product

- Does $\Phi(\mathbf{a})^T \Phi(\mathbf{b})$ look familiar?

$$\Phi(\mathbf{a})^T \Phi(\mathbf{b}) = 1 + 2 \sum_{i=1}^d a_i b_i + \sum_{i=1}^d a_i^2 b_i^2 + \sum_{i=1}^d \sum_{j=i+1}^d 2a_i a_j b_i b_j$$

- Try this: $(\mathbf{a}^T \mathbf{b} + 1)^2$

$$\begin{aligned} (\mathbf{a}^T \mathbf{b} + 1)^2 &= (\mathbf{a}^T \mathbf{b})^2 + 2\mathbf{a}^T \mathbf{b} + 1 \\ &= \left(\sum_{i=1}^d a_i b_i \right)^2 + 2 \sum_{i=1}^d a_i b_i + 1 \\ &= \sum_{i=1}^d \sum_{j=1}^d a_i b_i a_j b_j + 2 \sum_{i=1}^d a_i b_i + 1 \\ &= \sum_{i=1}^d a_i^2 b_i^2 + 2 \sum_{i=1}^d \sum_{j=i+1}^d a_i a_j b_i b_j + 2 \sum_{i=1}^d a_i b_i + 1 \end{aligned}$$

Quadratic Dot Product

- Does $\Phi(\mathbf{a})^T \Phi(\mathbf{b})$ look familiar?

$$\Phi(\mathbf{a})^T \Phi(\mathbf{b}) = 1 + 2 \sum_{i=1}^d a_i b_i + \sum_{i=1}^d a_i^2 b_i^2 + \sum_{i=1}^d \sum_{j=i+1}^d 2a_i a_j b_i b_j$$

- Try this: $(\mathbf{a}^T \mathbf{b} + 1)^2$

$$\begin{aligned} (\mathbf{a}^T \mathbf{b} + 1)^2 &= (\mathbf{a}^T \mathbf{b})^2 + 2\mathbf{a}^T \mathbf{b} + 1 \\ &= \left(\sum_{i=1}^d a_i b_i \right)^2 + 2 \sum_{i=1}^d a_i b_i + 1 \\ &= \sum_{i=1}^d \sum_{j=1}^d a_i b_i a_j b_j + 2 \sum_{i=1}^d a_i b_i + 1 \\ &= \sum_{i=1}^d a_i^2 b_i^2 + 2 \sum_{i=1}^d \sum_{j=i+1}^d a_i a_j b_i b_j + 2 \sum_{i=1}^d a_i b_i + 1 \end{aligned}$$

- They're the same! And this is only $O(d)$ to compute!

Dual SVM: Non-linear Transformation

$$\begin{aligned} \max_{\alpha \in \mathbb{R}^N} \quad & \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{m=1}^N \sum_{n=1}^N \alpha_n \alpha_m y_n y_m \Phi(\mathbf{x}_n)^T \Phi(\mathbf{x}_m) \\ \text{subject to} \quad & \sum_{n=1}^N y_n \alpha_n = 0, \alpha_n \geq 0, \forall n \end{aligned}$$

$$\mathbf{w}^* = \sum_{n=1}^N y_n \alpha_n^* \Phi(\mathbf{x}_n)$$

- Need to prepare a matrix Q , $Q_{nm} = y_n y_m \Phi(\mathbf{x}_n)^T \Phi(\mathbf{x}_m)$
- Cost?
 - We must do $N^2/2$ dot products to get this matrix ready.
 - Each dot product requires d additions and multiplications.

Higher Order Polynomials

	$\Phi(\mathbf{x})$	Cost	100dim
Quadratic	$d^2/2$ terms	$d^2 N^2/4$	$2.5kN^2$
Cubic	$d^3/6$ terms	$d^3 N^2/12$	$83kN^2$
Quartic	$d^4/24$ terms	$d^4 N^2/48$	$1.96mN^2$
	$\Phi(\mathbf{a})^T \Phi(\mathbf{b})$	Cost	100dim
Quadratic	$(\mathbf{a}^T \mathbf{b} + 1)^2$	$dN^2/2$	$50N^2$
Cubic	$(\mathbf{a}^T \mathbf{b} + 1)^3$	$dN^2/2$	$50N^2$
Quartic	$(\mathbf{a}^T \mathbf{b} + 1)^4$	$dN^2/2$	$50N^2$

Dual SVM with Quintic basis functions

$$\max_{\alpha \in \mathbb{R}^N} \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{m=1}^N \sum_{n=1}^N \alpha_n \alpha_m y_n y_m \underbrace{\Phi(\mathbf{x}_n)^T \Phi(\mathbf{x}_m)}_{(\mathbf{x}_n^T \mathbf{x}_m + 1)^5}$$

$$\text{subject to } \sum_{n=1}^N y_n \alpha_n = 0, \alpha_n \geq 0, \forall n$$

Classification:

$$\begin{aligned} g(\mathbf{x}) &= \text{sign}(\mathbf{w}^{*T} \Phi(\mathbf{x}) + b^*) = \text{sign} \left(\sum_{\alpha_n^* > 0} y_n \alpha_n^* \Phi(\mathbf{x}_n)^T \Phi(\mathbf{x}) + b^* \right) \\ &= \text{sign} \left(\sum_{\alpha_n^* > 0} y_n \alpha_n^* (\mathbf{x}_n^T \mathbf{x} + 1)^5 + b^* \right) \end{aligned}$$

Dual SVM with general kernel functions

$$\begin{aligned} \max_{\boldsymbol{\alpha} \in \mathbb{R}^N} \quad & \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{m=1}^N \sum_{n=1}^N \alpha_n \alpha_m y_n y_m K(\mathbf{x}_n, \mathbf{x}_m) \\ \text{subject to} \quad & \sum_{n=1}^N y_n \alpha_n = 0, \alpha_n \geq 0, \forall n \end{aligned}$$

Classification:

$$\begin{aligned} g(\mathbf{x}) &= \text{sign}(\mathbf{w}^{*T} \boldsymbol{\Phi}(\mathbf{x}) + b^*) = \text{sign} \left(\sum_{\alpha_n^* > 0} y_n \alpha_n^* \boldsymbol{\Phi}(\mathbf{x}_n)^T \boldsymbol{\Phi}(\mathbf{x}) + b^* \right) \\ &= \text{sign} \left(\sum_{\alpha_n^* > 0} y_n \alpha_n^* K(\mathbf{x}_n, \mathbf{x}_m) + b^* \right) \end{aligned}$$

Kernel Tricks

- Replacing dot product with a kernel function

Kernel Tricks

- Replacing dot product with a kernel function
- Not all functions are kernel functions!
 - Need to be decomposable $K(\mathbf{a}, \mathbf{b}) = \Phi(\mathbf{a})^T \Phi(\mathbf{b})$
 - Could $K(\mathbf{a}, \mathbf{b}) = (\mathbf{a} - \mathbf{b})^3$ be a kernel function?
 - Could $K(\mathbf{a}, \mathbf{b}) = (\mathbf{a} - \mathbf{b})^4 - (\mathbf{a} + \mathbf{b})^2$ be a kernel function?

Kernel Tricks

- Replacing dot product with a kernel function
- Not all functions are kernel functions!
 - Need to be decomposable $K(\mathbf{a}, \mathbf{b}) = \Phi(\mathbf{a})^T \Phi(\mathbf{b})$
 - Could $K(\mathbf{a}, \mathbf{b}) = (\mathbf{a} - \mathbf{b})^3$ be a kernel function?
 - Could $K(\mathbf{a}, \mathbf{b}) = (\mathbf{a} - \mathbf{b})^4 - (\mathbf{a} + \mathbf{b})^2$ be a kernel function?
- Mercer's condition
 - To expand Kernel function $K(\mathbf{a}, \mathbf{b})$ into a dot product, i.e., $K(\mathbf{a}, \mathbf{b}) = \Phi(\mathbf{a})^T \Phi(\mathbf{b})$, $K(\mathbf{a}, \mathbf{b})$ has to be positive semi-definite function.


Kernel Tricks

- Replacing dot product with a kernel function
- Not all functions are kernel functions!
 - Need to be decomposable $K(\mathbf{a}, \mathbf{b}) = \Phi(\mathbf{a})^T \Phi(\mathbf{b})$
 - Could $K(\mathbf{a}, \mathbf{b}) = (\mathbf{a} - \mathbf{b})^3$ be a kernel function?
 - Could $K(\mathbf{a}, \mathbf{b}) = (\mathbf{a} - \mathbf{b})^4 - (\mathbf{a} + \mathbf{b})^2$ be a kernel function?
- Mercer's condition
 - To expand Kernel function $K(\mathbf{a}, \mathbf{b})$ into a dot product, i.e., $K(\mathbf{a}, \mathbf{b}) = \Phi(\mathbf{a})^T \Phi(\mathbf{b})$, $K(\mathbf{a}, \mathbf{b})$ has to be positive semi-definite function.
 - kernel matrix K is always symmetric PSD for any given $\mathbf{x}_1, \dots, \mathbf{x}_N$.


Kernel Design: expression kernel

- mRNA expression data:

- Each matrix entry is an mRNA expression measurement.
- Each column is an experiment.
- Each row corresponds to a gene.



Kernel Design: expression kernel



- mRNA expression data:
 - Each matrix entry is an mRNA expression measurement.
 - Each column is an experiment.
 - Each row corresponds to a gene.

- Similar or dissimilar




Similar



Dissimilar

Kernel Design: expression kernel



- mRNA expression data:
 - Each matrix entry is an mRNA expression measurement.
 - Each column is an experiment.
 - Each row corresponds to a gene.
- Similar or dissimilar
 - Similar
 - Dissimilar
- Kernel

$$K(x, y) = \frac{\sum_i x_i y_i}{\sqrt{\sum_i x_i x_i} \sqrt{\sum_i y_i y_i}}$$

Kernel Design: sequence kernel

- Work with non-vectorial data
- Scalar product on a pair of variable-length, discrete strings?

>ICYA _MANSE

GDIFYPGYCPDVKPVNDFDLSAFAGAWHEIAKLPLENENQKGCTIAEYKY
DGKKASVYNSFVSNVKEYMEGDLEIAPDAKYTKQKGKYVMTFKFGQRVVN
LVPWVLATDYKNYAINYMENSHPDKKAHSIHAWILSKSKVLEGNTKEVVD
NVLKTFSHLIDASKFISNDFSEAACQYSTTYSLTGPDRH

>LACB_BOVIN

MKCLLLALALTCTGAQALIVTQTMKGLDIQKVAGTWYSLAMAASDI SLLDA
QSAPLRVYVEELKPTPEGDLEILLQKWENGECQAQKKI IAEKTKIPAVFKI
DALNENKVLVLDTDYKKYLLFCMENSAEPEQSLACQCLVRTPEVDDEALE
KFDKALKALPMHIRLSFNPTQLEEQCHI

Commonly Used SVM Kernel Functions


- $K(\mathbf{a}, \mathbf{b}) = (\alpha \cdot \mathbf{a}^T \mathbf{b} + \beta)^Q$ is an example of a SVM kernel function.
- Beyond polynomials there are other very high dimensional basis functions that can be made practical by finding the right Kernel Function
 - Radial-basis style kernel (RBF)/Gaussian kernel function

$$K(\mathbf{a}, \mathbf{b}) = \exp(-\gamma \|\mathbf{a} - \mathbf{b}\|^2)$$


- Sigmoid functions

2nd Order Polynomial Kernel


$$K(\mathbf{a}, \mathbf{b}) = (\alpha \cdot \mathbf{a}^T \mathbf{b} + \beta)^2$$



$$(1 + 0.001\mathbf{x}^T \mathbf{x}')^2$$




$$1 + (\mathbf{x}^T \mathbf{x}') + (\mathbf{x}^T \mathbf{x}')^2$$




$$(1 + 1000\mathbf{x}^T \mathbf{x}')^2$$

Gaussian Kernels


$$K(\mathbf{a}, \mathbf{b}) = \exp(-\gamma \|\mathbf{a} - \mathbf{b}\|^2)$$



$$\exp(-1\|\mathbf{x} - \mathbf{x}'\|^2)$$




$$\exp(-10\|\mathbf{x} - \mathbf{x}'\|^2)$$



$$\exp(-100\|\mathbf{x} - \mathbf{x}'\|^2)$$


Gaussian Kernels

$$K(\mathbf{a}, \mathbf{b}) = \exp(-\gamma \|\mathbf{a} - \mathbf{b}\|^2)$$




When γ is large, we clearly see that even the protection of a large margin cannot suppress overfitting. However, for a reasonably small γ , the sophisticated boundary discovered by SVM with the Gaussian-RBF kernel looks quite good.

Gaussian Kernels



(a) linear classifier



(b) Gaussian-RBF kernel

For (a) a noisy data set that linear classifier appears to work quite well, (b) using the Gaussian-RBF kernel with the hard-margin SVM leads to overfitting.

From hard-margin to soft-margin

- When there are outliers, hard margin SVM + Gaussian-RBF kernel result in an unnecessarily complicated decision boundary that overfits the training noise.

From hard-margin to soft-margin

- When there are outliers, hard margin SVM + Gaussian-RBF kernel result in an unnecessarily complicated decision boundary that overfits the training noise.
- Remedy: a soft formulation that allows small violation of the margins or even some classification errors.

From hard-margin to soft-margin

- When there are outliers, hard margin SVM + Gaussian-RBF kernel result in an unnecessarily complicated decision boundary that overfits the training noise.
- Remedy: a soft formulation that allows small violation of the margins or even some classification errors.
- Soft-margin: margin violation $\varepsilon_n \geq 0$ for each data point (\mathbf{x}_n, y_n) and require that

$$y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 - \varepsilon_n$$

- ε_n captures by how much (\mathbf{x}_n, y_n) fails to be separated.

Soft-Margin SVM

We modify the hard-margin SVM to the soft-margin SVM by allowing margin violations but adding a penalty term to discourage large violations:

$$\min_{b, \mathbf{w}, \boldsymbol{\varepsilon}} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{n=1}^N \varepsilon_n$$

subject to: $y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 - \varepsilon_n$ for $n = 1, \dots, N$

$\varepsilon_n \geq 0$, for $n = 1, \dots, N$

The meaning of C ?

Soft-Margin SVM

We modify the hard-margin SVM to the soft-margin SVM by allowing margin violations but adding a penalty term to discourage large violations:

$$\min_{b, \mathbf{w}, \epsilon} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{n=1}^N \epsilon_n$$


subject to: $y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 - \epsilon_n$ for $n = 1, \dots, N$

$\epsilon_n \geq 0$, for $n = 1, \dots, N$


The meaning of C ?

- When C is large, it means we care more about violating the margin, which gets us closer to the hard-margin SVM.
- When C is small, on the other hand, we care less about violating the margin.

Soft Margin Example



(a) $C = 1$




(b) $C = 500$

Soft Margin and Hard Margin

$$\min_{b, \mathbf{w}, \varepsilon} \underbrace{\frac{1}{2} \mathbf{w}^T \mathbf{w}}_{\text{margin}} + C \underbrace{\sum_{n=1}^N \varepsilon_n}_{\text{error tolerance}}$$

subject to: $y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 - \varepsilon_n, \varepsilon_n \geq 0, \forall N$



The Hinge Loss

- The trade-off sounds very similar, right?

The Hinge Loss

- The trade-off sounds very similar, right?
- We have $\varepsilon_n \geq 0$, and that
$$y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 - \varepsilon_n \Rightarrow \varepsilon_n \geq 1 - y_n(\mathbf{w}^T \mathbf{x}_n + b)$$

The Hinge Loss

- The trade-off sounds very similar, right?
- We have $\varepsilon_n \geq 0$, and that
$$y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 - \varepsilon_n \Rightarrow \varepsilon_n \geq 1 - y_n(\mathbf{w}^T \mathbf{x}_n + b)$$
- The SVM loss (aka. Hinge Loss) function

$$E_{\text{SVM}}(b, \mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \max(1 - y_n(\mathbf{w}^T \mathbf{x}_n + b), 0)$$

The Hinge Loss

- The trade-off sounds very similar, right?
- We have $\varepsilon_n \geq 0$, and that
$$y_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 - \varepsilon_n \Rightarrow \varepsilon_n \geq 1 - y_n(\mathbf{w}^T \mathbf{x}_n + b)$$
- The SVM loss (aka. Hinge Loss) function

$$E_{\text{SVM}}(b, \mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \max(1 - y_n(\mathbf{w}^T \mathbf{x}_n + b), 0)$$

- The soft-margin SVM can be re-written as the following optimization problem:

$$\min_{b, \mathbf{w}} E_{\text{SVM}}(b, \mathbf{w}) + \lambda \mathbf{w}^T \mathbf{w}$$

Dual Soft-Margin SVM

$$\begin{aligned} & \max_{\boldsymbol{\alpha} \in \mathbb{R}^N} \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{m=1}^N \sum_{n=1}^N \alpha_n \alpha_m y_n y_m \mathbf{x}_n^T \mathbf{x}_m \\ & \text{subject to } \sum_{n=1}^N y_n \alpha_n = 0, 0 \leq \alpha_n \leq C, \forall n \end{aligned}$$

$$\mathbf{w}^* = \sum_{n=1}^N y_n \alpha_n^* \mathbf{x}_n$$

Summary of Dual SVM

- Deliver a large-margin hyperplane, and in so doing it can control the effective model complexity.
- Deal with high- or infinite-dimensional transforms using the kernel trick.
- Express the final hypothesis $g(\mathbf{x})$ using only a few support vectors, their corresponding dual variables (Lagrange multipliers), and the kernel.
- Control the sensitivity to outliers and regularize the solution through setting C appropriately.

Support Vector Machine

